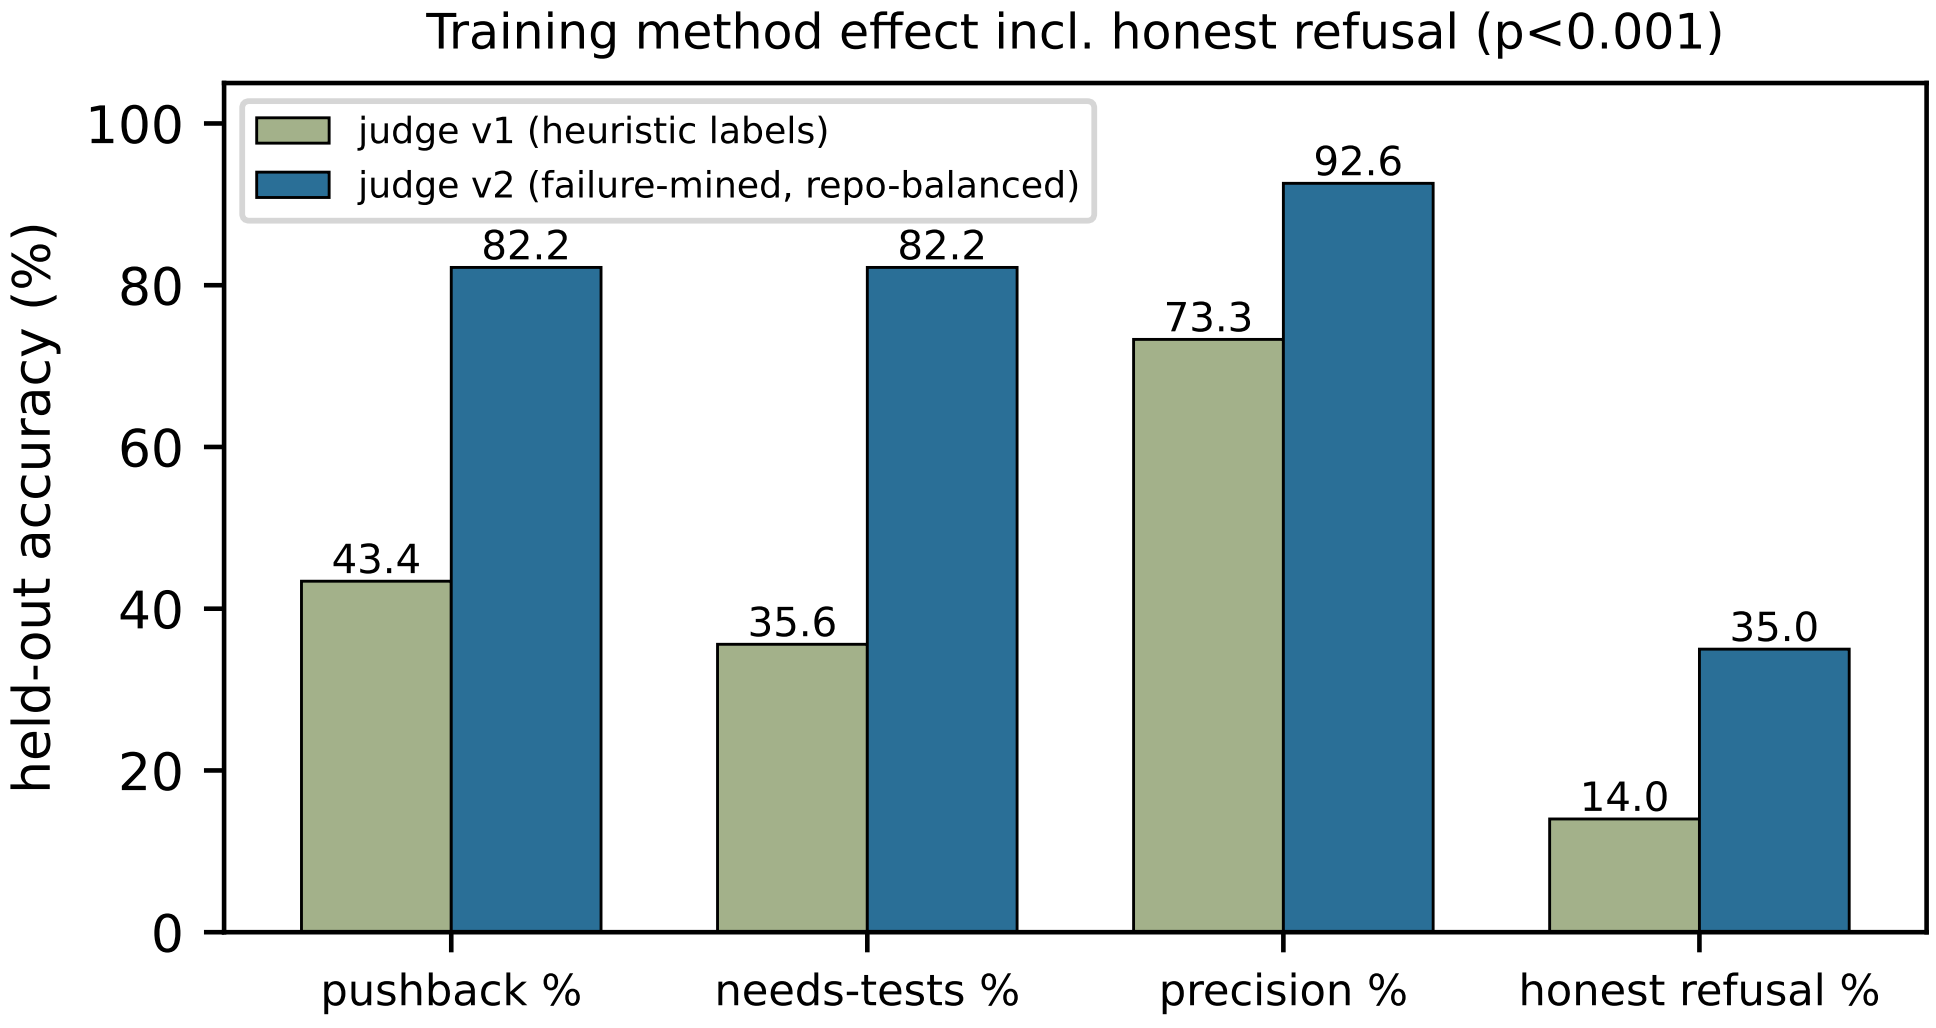
Training method effect incl. honest refusal (p<0.001)
judge v1 (heuristic labels)
judge v2 (failure-mined, repo-balanced)
held-out accuracy (%)
0
20
40
60
80
100
43.4
82.2
35.6
82.2
73.3
92.6
14.0
35.0
pushback %
needs-tests %
precision %
honest refusal %